

HyperLogistics: Explainable Rerouting & Safety Justification

Domain Adaptation Group Report

Tony Nguyen

Joel Vinas

Daniel Evans

March 16, 2026

1. Project Description & Domain Task

HyperLogistics is a Snowflake-native supply chain resilience system that bridges the prediction-action gap for middle-mile logistics. It uses *ReMindRAG* for knowledge-guided retrieval and *SRSNet* for adaptive forecasting, ensuring autonomous rerouting strategies grounded in safety and compliance.

The core **Domain Task** is **Explainable Rerouting & Safety Justification**. This involves generating constraint-compliant rerouting justifications for middle-mile logistics dispatchers when real-time disruptions (e.g., weather alerts or accident blackspots) are present. The expected output is a structured rationale that proposes a reroute, cites the disruption, and confirms compliance with strict physical constraints (e.g., DOT bridge weight/height limits).

2. Instruction Dataset Description

To teach the model domain-specific constraints (e.g., DOT bridge limits) and logistics jargon (e.g., “heavy haul”, “LTL”, “bobtail”), we generated a synthetic dataset. A pipeline (`generate_dataset.py`) simulated Area Manager queries regarding real-time disruptions, programmatically combining constraints from `SILVER.BRIDGE_INVENTORY_GEO` and simulated events from `SILVER.WEATHER_ALERTS`. This yielded 100 query-response pairs that establish the exact reasoning patterns the model must follow to safely veto or approve a route. The dataset is saved as `instruction_dataset.json`.

3. Adaptation Method Used

We employed a **multi-strategy** adaptation approach combining Parameter-Efficient Fine-Tuning (PEFT) with advanced prompt engineering:

- **PEFT Fine-Tuning (QLoRA):** Using the `microsoft/phi-2` (2.7B) base model, we applied 4-bit NF4 quantization via `bitsandbytes`. The LoRA config targeted `q_proj` and `v_proj` layers (Rank=16, Alpha=32) with a `paged_adamw_8bit` optimizer for 3 epochs. The lightweight PEFT adapter is loaded on top of the base model at inference time.
- **Advanced Prompt Adaptation:** Using `prompt_adaptation.py`, we implemented *Few-Shot* (injecting similar validated examples), *Self-Consistent Chain-of-Thought (SC-CoT)* (majority vote across 3 reasoning chains), and *ReAct* (interleaved Thought/Action steps for grounded reasoning).

4. System Integration Description

The domain-adapted model is integrated directly into the broader HyperLogistics architecture, a 4-layer Snowflake-native system:

1. **Data Perception:** Raw data arrives via Snowpipe (US Accidents, NOAA Weather) and Snowflake Stages (NBI Bridge constraints).
2. **Intelligence & Forecasting:** ReMindRAG acts as a reasoning agent, and SRSNet provides 4–8h risk propagation forecasts before route generation.
3. **Validation & Safety:** The Consensus Planning Protocol (CPP) utilizes the domain-adapted PEFT model as a Safety Validation Agent to negotiate routes and enforce DOT bridge limits deterministically using Spatial SQL. No LLM reasoning can override a physical constraint veto.
4. **Application:** A Streamlit Dashboard (`demo_dashboard.py`) processes dispatcher queries via the adapted model to produce structured, constraint-compliant justifications.

5. Evaluation Results

We evaluated the baseline Phi-2 model against our PEFT-adapted model using a custom 15-query dataset featuring metamorphic testing pairs (invariance, monotonicity, and symmetry).

- **Standard Metrics:** The **Baseline Accuracy was 40.0%**, while the **Adapted PEFT Accuracy reached 95.0%**. The un-tuned SC-CoT strategy achieved 93.0%. Domain Relevance, Grounding, and Response Clarity all saw massive improvements (from 33% to near 100%).
- **Local Hardware Evaluation (RTX 3060):** A real-hardware test (4-bit inference) confirmed structural reasoning learned through fine-tuning, but identified monotonicity failures. The primary failure mode was hallucination—the model frequently invented bridge constraints not present in the evidence, then correctly vetoed based on its own made-up data.

6. Impact on Project Performance

The baseline model achieved only 40% accuracy with generic, non-grounded responses. Our best adapted strategy (PEFT) achieved 95% accuracy with fully constraint-aware, structured reasoning—a **+55 percentage point improvement**. This demonstrates that the adaptation transformed the application from a generic text generator into a compliance-focused decision agent. The reasoning structures learned via CoT and fine-tuning are sound, highlighting that bridging the prediction-action gap requires explicit domain adaptation for deterministic safety rules.